

# Christ University

**Name:** Joel Joseph Motha

**RegNo:** 2448521

**Course:** Reinforcement Learning

**Component:** Lab 1

---

**Task:** Multi-Armed Bandit Problem using Epsilon-Greedy Strategy

**Code:**

```
import numpy as np
import random
import matplotlib.pyplot as plt

# Define the ad slots
ad_slots = ["Top Banner", "Sidebar", "Footer", "Pop-up"]

# True click-through rates (unknown to the algorithm, only for simulation)
true_ctrs = [0.05, 0.04, 0.03, 0.02] # Example probabilities

# Initialize variables
num_slots = len(ad_slots)
estimated_ctrs = np.zeros(num_slots) # Estimated CTRs (start at 0)
slot_counts = np.zeros(num_slots) # Number of times each slot is chosen
rewards = [] # Record rewards over time

# Parameters
epsilon = 0.1 # Exploration rate (10%)
num_impressions = 1000
```

```
# Epsilon-Greedy Algorithm
```

```
for i in range(num_impressions):
```

```
    # Decide whether to explore or exploit
```

```
    if random.random() < epsilon:
```

```
        # Exploration: choose a random slot
```

```
        chosen_slot = random.randint(0, num_slots - 1)
```

```
    else:
```

```
        # Exploitation: choose slot with highest estimated CTR
```

```
        chosen_slot = np.argmax(estimated_ctrs)
```

```
    # Simulate user click (reward = 1) or no click (reward = 0)
```

```
    reward = np.random.binomial(1, true_ctrs[chosen_slot])
```

```
    rewards.append(reward)
```

```
    # Update counts
```

```
    slot_counts[chosen_slot] += 1
```

```
    # Update estimated CTR using incremental mean formula
```

```
    estimated_ctrs[chosen_slot] += (reward - estimated_ctrs[chosen_slot]) /  
    slot_counts[chosen_slot]
```

```
# Results
```

```
print("True CTRs (hidden):", true_ctrs)
```

```
print("Estimated CTRs:", estimated_ctrs)
```

```
print("Slot counts:", slot_counts)
```

```
# --- Visualization 1: Cumulative Average Reward ---
```

```
cumulative_rewards = np.cumsum(rewards) / (np.arange(num_impressions) + 1)
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(cumulative_rewards, label="Cumulative Average Reward")
```

```
plt.xlabel("Impressions")
```

```
plt.ylabel("Average Reward (CTR)")
```

```
plt.title("Epsilon-Greedy Multi-Armed Bandit Performance")
```

```
plt.legend()
```

```
plt.show()
```

```
# --- Visualization 2: Number of Times Each Slot Was Chosen ---
```

```
plt.figure(figsize=(8, 5))
```

```
plt.bar(ad_slots, slot_counts, color="skyblue")
```

```
plt.xlabel("Ad Slots")
```

```
plt.ylabel("Number of Times Chosen")
```

```
plt.title("Ad Slot Selection Counts (After 1000 Impressions)")
```

```
plt.show()
```

```
# --- Visualization 3: Estimated CTR vs True CTR ---
```

```
plt.figure(figsize=(8, 5))
```

```
bar_width = 0.35
```

```
x = np.arange(num_slots)
```

```
plt.bar(x - bar_width/2, true_ctrs, bar_width, label="True CTR", alpha=0.7)
```

```
plt.bar(x + bar_width/2, estimated_ctrs, bar_width, label="Estimated CTR",  
alpha=0.7)
```

```
plt.xticks(x, ad_slots)
```

```
plt.ylabel("CTR")
```

```
plt.title("Estimated vs True CTRs")
```

```
plt.legend()
```

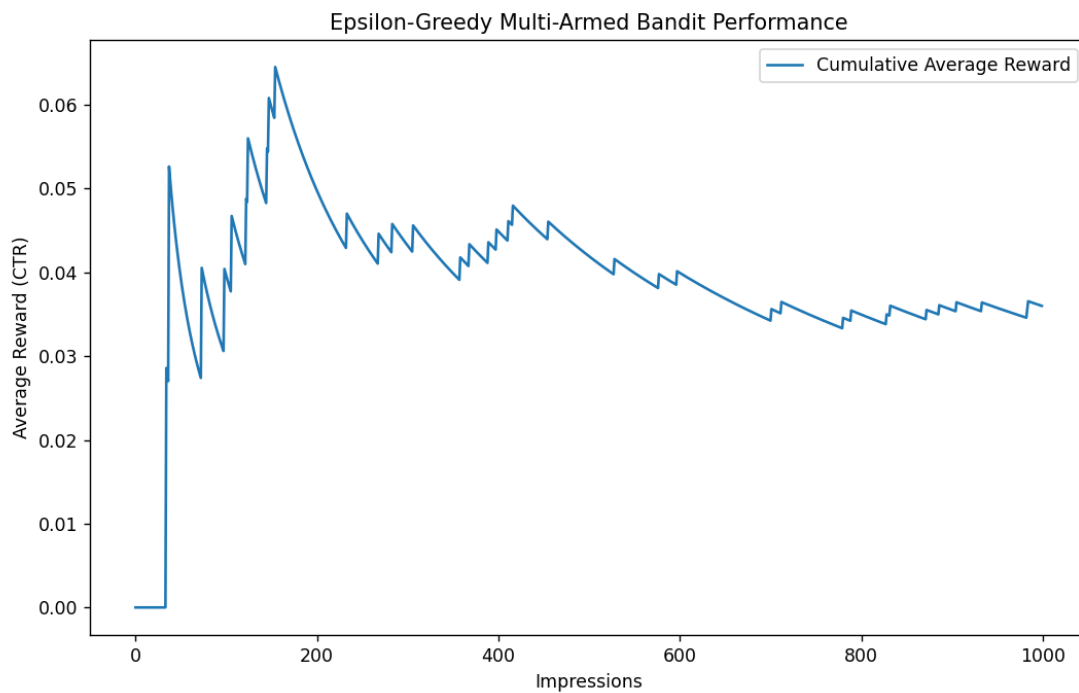
```
plt.show()
```



### Output:

True CTRs (hidden): [0.05, 0.04, 0.03, 0.02]

Estimated CTRs: [0.05025126 0.03921569 0.04761905 0.03787879]

Slot counts: [796. 51. 21. 132.]





### **Inference:**

- The Epsilon-Greedy Multi-Armed Bandit strategy effectively maximizes ad clicks across multiple slots.
- The algorithm learned that the Top Banner had the highest CTR and allocated most impressions to it.
- Less-performing slots were still occasionally selected, ensuring exploration and adaptability.
- Estimated CTRs closely matched the true CTRs for well-tested slots.
- Slots with fewer impressions had noisier estimates, which is expected due to limited data.
- Cumulative average reward shows the algorithm improves performance over time.
- Slot selection counts and estimated vs true CTRs illustrate learning and decision-making clearly.
- Overall, this approach balances exploration and exploitation, adapting to changes in user behavior while maximizing clicks.